

Inter Process Communication(IPC)

- ❑ **IPC between process on different host**

- ❑ **Sockets**

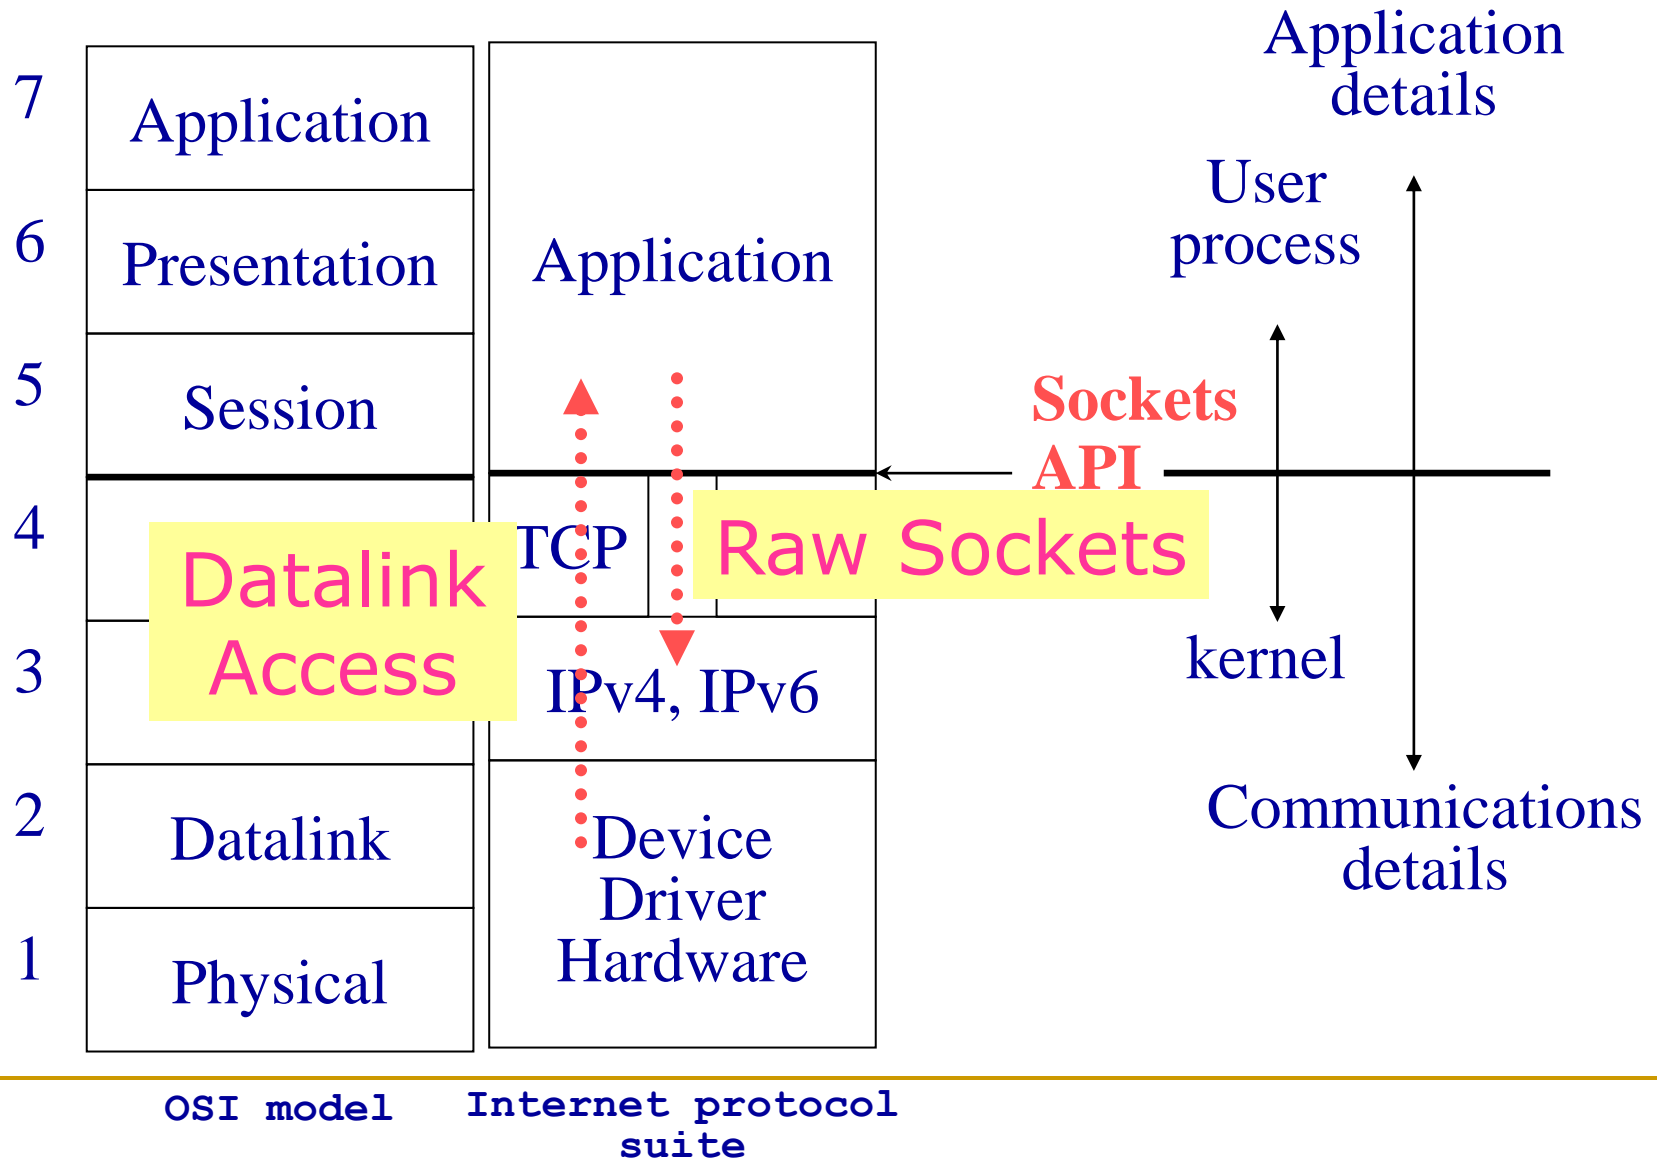
- ❑ **TCP Socket (Stream Socket)**

- ❑ **UDP Socket (Datagram Socket)**

- ❑ **SCTP Sockets**

- ❑ **Raw Socket (bypass Transport Layer)**

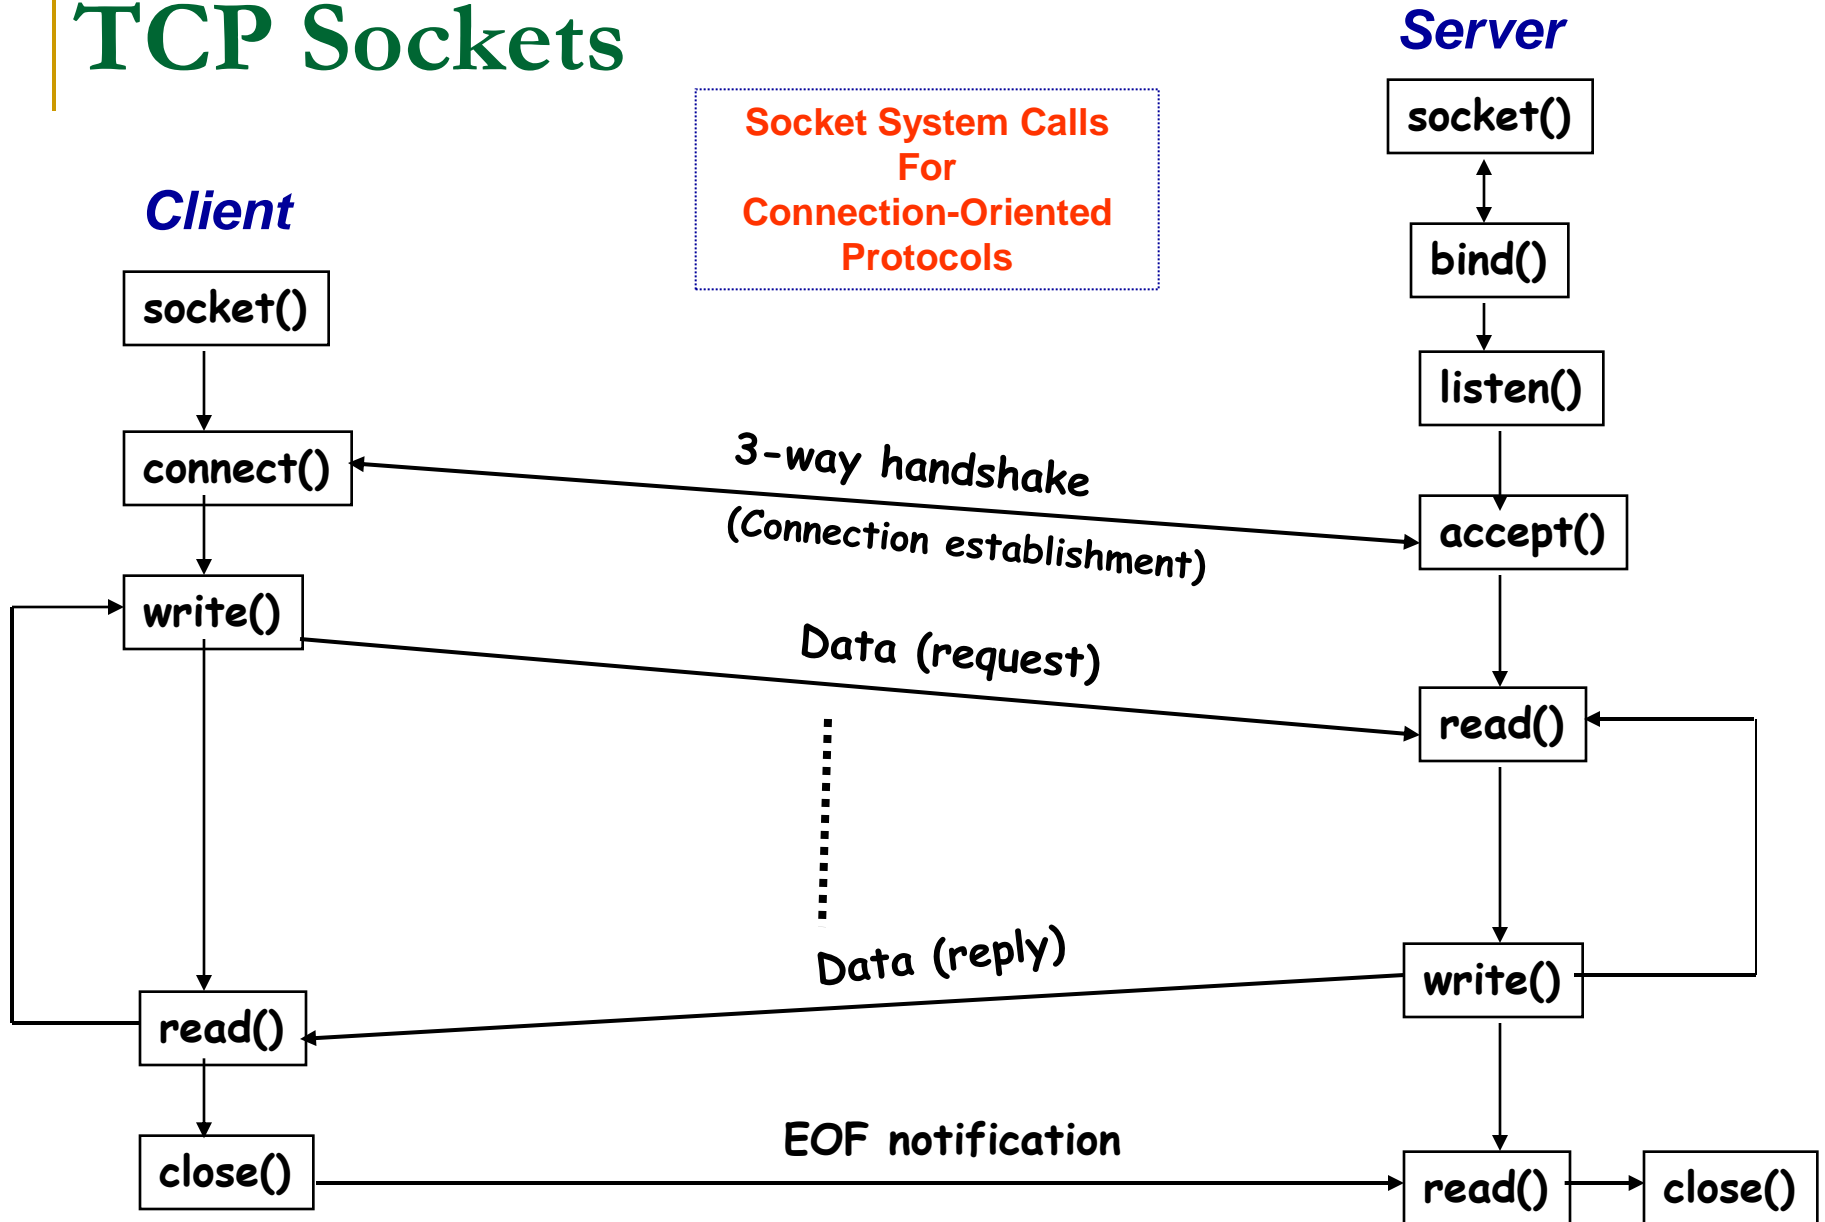
Why do socket provides interface from upper three layers of OSI into transport layer ??



Client \leftrightarrow Server

- ❑ A *server* is a process - not a machine !
 - ❑ A server waits for a request from a client
 - ❑ A client is a process that sends a request to an existing server and (usually) waits for a reply.
 - ❑ In a network, the client/server model provides a convenient way to interconnect programs that are distributed efficiently across different locations
-

TCP Sockets



Socket System Calls

- Creates a network plug point that enables the client/server to communicate

```
#include <sys/types.h>
#include <sys/socket.h>

int socket (int family, int type, int protocol);
```

Socket descriptor
-1 if fails

AF_INET Internet protocols (IPv4)
AF_LOCAL Unix domain protocols
AF_ROUTE Routing sockets
AF_NS Xerox NS protocol
AF_INET6 IPv6 Protocol

Usually 0 is used.
Defined in <netinet/in.h>
IPPROTO_XXX

	AF_UNIX	AF_INET
SOCK_STREAM	Yes	TCP
SOCK_DGRAM	Yes	UDP
SOCK_RAW		IP

SOCK_STREAM	Stream Socket
SOCK_DGRAM	Datagram Socket
SOCK_RAW	Raw Socket
SOCK_SEQPACKET	Sequenced Packet Socket

Bind System Calls

- ❑ Allows a server to specify the IP address and port number associated with a socket

```
#include <sys/types.h>
#include <sys/socket.h>

int bind (int sockfd, const struct sockaddr *myaddr, int addrlen);
```

Pointer to a protocol-specific address



- ❑ Assign local protocol address to socket
 - ❑ Servers register their addresses with the system
 - ❑ Both connection oriented & connectionless servers need to do this before accepting client requests.
 - ❑ Client can register a specific address for itself
 - ❑ connectionless clients need this to assure that a specific address is assigned, so that it can receive response to its request

Connect System Calls

- ❑ Enables a client to connect to a server i.e. initiate a connection to a remote host

```
#include <sys/types.h>
```

```
#include <sys/socket.h>
```

```
int connect (int sockfd, struct sockaddr *servaddr, int addrlen);
```

- ❑ for connection oriented protocols, like TCP, results in actual connection establishment (3-wayhandshake)
 - ❑ Client process connects a socket descriptor to establish a connection with the server
 - ❑ Specific communication parameters like buffer size, size of data to be exchanged without ACK, etc, is agreed
 - ❑ System call does not return until connection is established or an error occurs
 - ❑ `bind` is not required before `connect`
- ❑ Connect can also used for connectionless client

Listen system calls

- ❑ allows the server to specify that a socket can be used to accept connections

```
int listen (int sockfd, int backlog);
```

- ❑ used by connection oriented server, to indicate that it is willing to receive connections
 - ❑ *backlog* specifies the number of requests that can be queued by the system, while it waits for the server to execute the `accept` system call

Accept system call

- allows a server to wait till a new connection request arrives

```
#include <sys/types.h>
#include <sys/socket.h>

int accept (int sockfd, struct sockaddr *cliaddr, int *addrlen);
```

- after `listen`, `accept` waits for an actual connection request from a client
 - takes the first connection request from queue, and creates another socket with the same properties as *sockfd*
 - if no requests are in queue, blocks until a connection request is received
 - `accept` automatically creates a new socket descriptor, assuming the server is a concurrent server

Read & Write system calls

- to read data from & write data to a TCP socket connection

```
#include <sys/types.h>
#include <sys/socket.h>

int read (int sockfd, char *buff, int nbytes);
int write (int sockfd, char *buff, int nbytes);

int readn (int sockfd, char *buff, int nbytes);
int writen (int sockfd, char *buff, int nbytes);
int writen (int sockfd, char *buff, int nbytes);
int readline (int sockfd, char *buff, int maxlen);
```

not built-in function

- all return the length of data that was read or written as the value of the function

Close system calls

- ❑ terminates any connection associated with a socket and releases the socket descriptor

```
int close (int sockfd);
```

- ❑ Normal UNIX call close is called to close a socket
 - ❑ normally unblocking even if, reliable transfer is required and system has still data or acknowledgements to be sent
 - ❑ Kernel still tries to handle those, although the function immediately returns

Getting IP address/port from socket

❑ **getpeername System Call**

- ❑ Get the IP/port of remote endpoint

```
#include <sys/types.h>
#include <sys/socket.h>

int getpeername(int sockfd, struct sockaddr *peername, int *addrlen );
```

❑ **getsockname System Call**

- ❑ Get the local IP/port bound to socket

```
#include <sys/types.h>
#include <sys/socket.h>

int getsockname(int sockfd, struct sockaddr *localaddr, int *addrlen );
```

Iterative server

```
int sockfd, newsockfd;

if ( (sockfd = socket ( ... )) < 0)
    err_sys ("socket error");
if (bind(sockfd, ... ) < 0)
    err_sys ("bond error");
if (listen(sockfd, 5) < 0)
    err_sys ("listen error");

for ( ; ; ) {
    newsocfd = accept (sockfd, ... );    /* blocking */
    if (newsockfd < 0)
        err_sys ("accept error");
    do_function();    /* process request */
    close (newsockfd);
}
```

Concurrent Server

```
int sockfd, newsockfd;

if ( (sockfd = socket ( ... )) < 0)
    err_sys ("socket error");
if (bind(sockfd, ... ) < 0)
    err_sys ("bond error");
if (listen(sockfd, 5) < 0)
    err_sys ("listen error");

for ( ; ; ) {
    newsockfd = accept (sockfd, ... );    /* blocking */
    if (newsockfd < 0)
        err_sys ("accept error");
    if (fork() == 0) {                    /* child process*/
        close (sockfd);
        do_function(newsockfd); /* process request */
        close (newsockfd);
        exit(0);
    }
    close (newsockfd);                    /* parent process */
}
```

Bind Example

```
#define PORT 989898
#define IPADD "172.24.2.4"
int sockfd;
struct sockaddr_in myaddr;
sockfd = socket(AF_INET,SOCK_STREAM,0);
    myaddr.sin_family = AF_INET;
    myaddr.sin_addr.s_addr = inet_addr( IPADD);
                                //or inet_addr(argv[1] ) for command line arg
    myaddr.sin_port = htons( PORT);    // or htons(atoi (argv[2]) ) for command line arg

bind(sockfd , (struct sockaddr*)&myaddr, sizeof(myaddr));
```

- ❑ during call to bind we can tell kernel to assign any available port
myaddr.sin_port = htons(0); //emhemeral port
- ❑ to tell kernel to choose IP address
myaddr.sin_addr.s_addr = htonl(INADDR_ANY); //wildcard address
- ❑ Clients typically don't care what port they are assigned

```
#define SERVER_PORT 898989
```

```
main {
```

```
struct sockaddr_in cliaddr,servaddr;
```

```
int sockfd,confd,clilen;
```

```
sockfd = socket(AF_INET , SOCK_STREAM , 0) ;
```

```
servaddr.sin_family = AF_INET;
```

```
servaddr.sin_port = htons(SERVER_PORT);
```

```
servaddr.sin_addr.s_addr =htonl( INADDR_ANY) ;
```

```
bind(sockfd, (struct sockaddr *)&servaddr,sizeof(servaddr));
```

```
listen(sockfd,10);
```

```
while(1) {
```

```
clilen=sizeof(cliaddr);
```

```
confd= accept(sockfd,(struct sockaddr *)&cliaddr,&clilen);
```

```
// /* printf("Client IP: %s\n", inet_ntoa(cliaddr.sin_addr)); */
```

```
// /* printf("Client Port: %hu\n", ntohs(cliaddr.sin_port)); */
```

```
read(connfd, buf, BUFFER_SIZE);
```

```
printf("Received:%s\n",buf);
```

```
write(confd, buf, BUFFER_SIZE);
```

```
close(confd);
```

```
}
```

```
return 0;
```

```
}
```

Iterative Server

\$/server

Note:

- Port number is hard-coded in program

-it can be taken as command line argument


```

#define SERVER_PORT 898989
main {
    struct sockaddr_in cliaddr,servaddr;
    int sockfd,confd,clilen;
sockfd = socket(AF_INET , SOCK_STREAM , 0) ;

    servaddr.sin_family = AF_INET;
servaddr.sin_port = htons(SERVER_PORT);
servaddr.sin_addr.s_addr =inet_addr(argv[1]) ; // or inet_pton()

connect(sockfd, (struct sockaddr *) &servaddr, sizeof(servaddr)) ;
    clilen = sizeof(cliaddr);

    // /* getsockname(sockfd, (struct sockaddr *) &cliaddr, &clilen); */
    // /* printf("Client socket has IP: %s\n", inet_ntoa(cliaddr.sin_addr)); */
    // /* printf("Client socket has Port: %hu\n", ntohs(cliaddr.sin_port)); */

    printf("Enter data:\n"); scanf("%s",buf) ;
write(sockfd, buf, BUFFER_SIZE);
read(sockfd, buf, BUFFER_SIZE);
printf("Received Data :%s\n",buf);
    close(sockfd);
    exit(0);
}

```

Client Program
\$/client 172.24.2.4

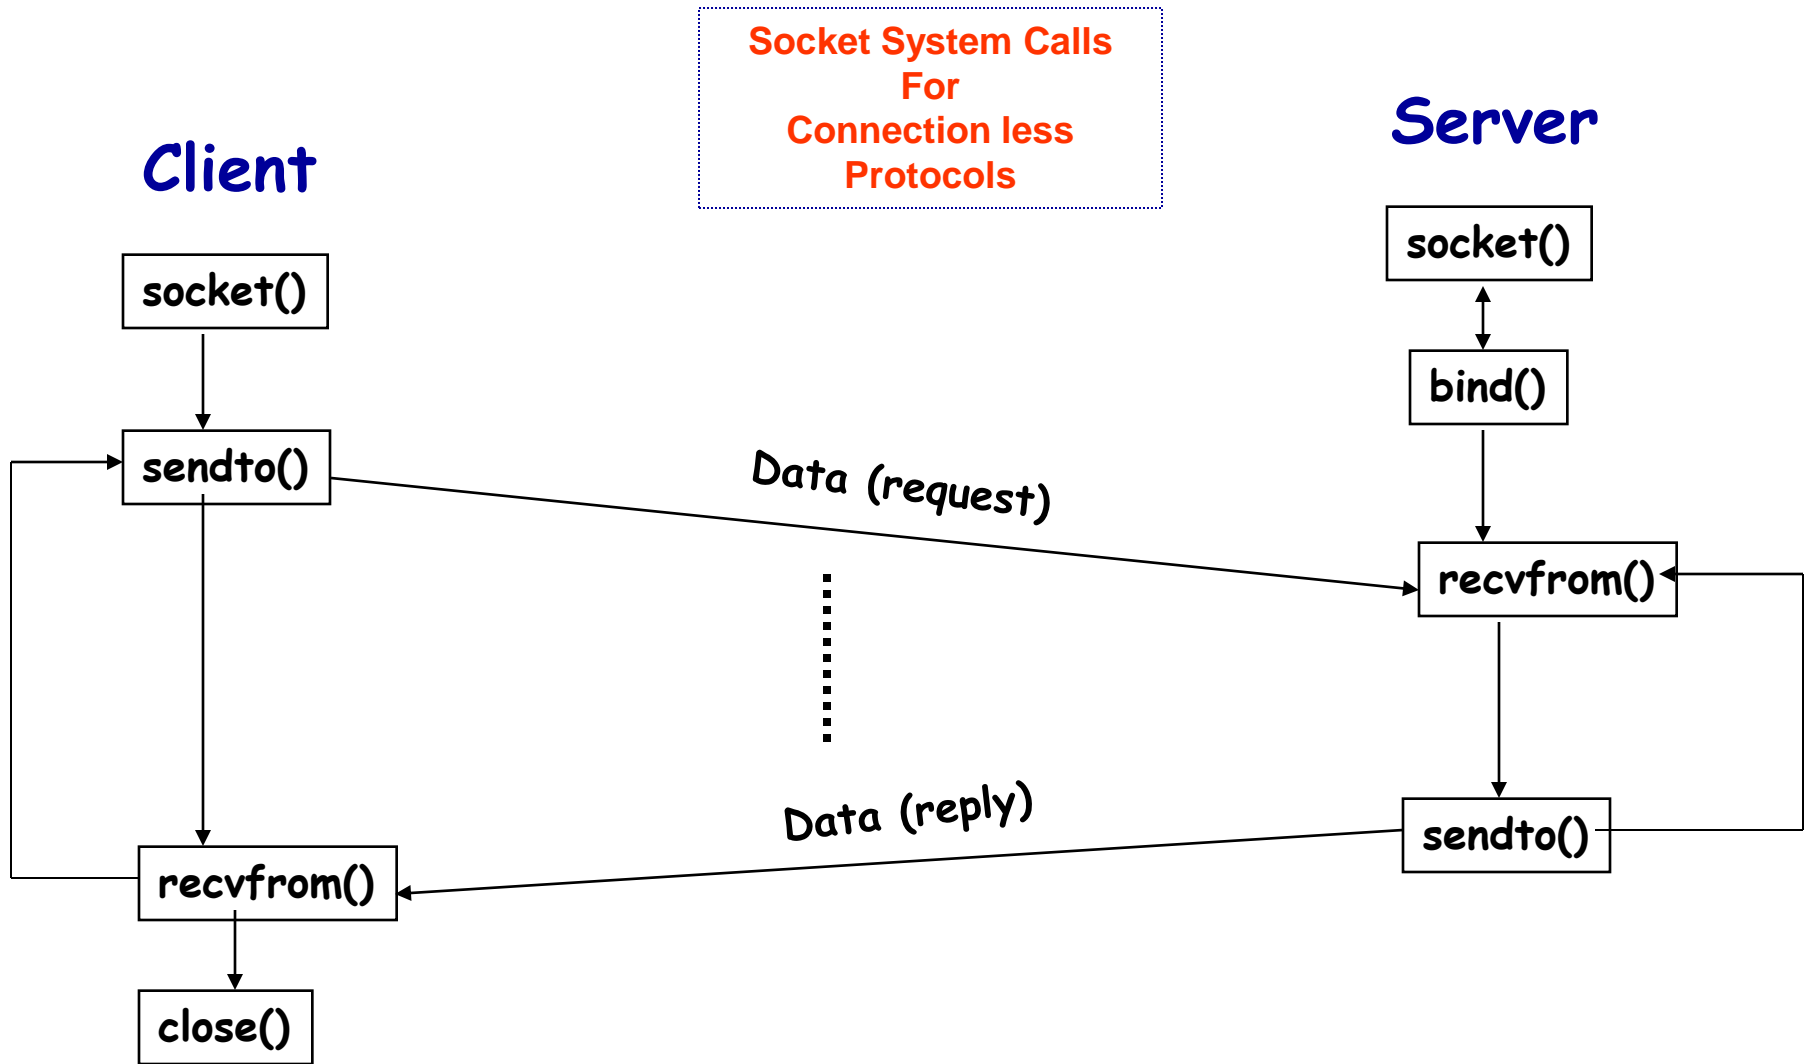
Note:

- Port number is hard-coded in program
- it can be taken as command line argument

concurrent server

```
main() {  
    sockfd = socket(...);  
    bind(sockfd,.....);  
    listen(sockfd,10);  
  
    for ( ; ; )  
    {  
        clien=sizeof(cliaddr);  
        confd= accept(sockfd,(struct sockaddr *)&cliaddr,&clilen);  
        printf("Client IP: %s\n", inet_ntoa(cliaddr.sin_addr));  
        printf("Client Port: %hu\n", ntohs(cliaddr.sin_port));  
  
        if( (pid = fork() ) ==0 ) /*child process handle the client request*/  
        { close(sockfd); bzero(buf,BUFFER_SIZE);  
          read(connfd, buf, BUFFER_SIZE);  
          printf("Received:%s\n",buf);  
          write(confd, buf, BUFFER_SIZE);  
          close(confd);  
          exit(0);  
        }  
        close(confd);  
    }  
}
```

UDP Sockets



Socket System Calls

- ❑ Creates a network plug point that enables the client/server to communicate

```
#include <sys/types.h>
#include <sys/socket.h>
int socket (int family, int type, int protocol);
```

Socket descriptor
-1 if fails

AF_INET

Internet protocols (IPv4)

Usually 0 is used.
Defined in <netinet/in.h>
IPPROTO_XXX

SOCK_DGRAM

Datagram Socket

Bind System Calls

- ❑ Allows a server to specify the IP address and port number associated with a socket

```
#include <sys/types.h>
#include <sys/socket.h>

int bind (int sockfd, const struct sockaddr *myaddr, int addrlen);
```

Pointer to a protocol-specific address



- ❑ Assign local protocol address to socket
 - ❑ Servers register their addresses with the system
 - ❑ Client can register a specific address for itself
 - ❑ connectionless clients need this to assure that a specific address is assigned, so that it can receive response to its request

Sendto System Calls

Either **zero**, or OR'ed with
MSG_OOB Out-of-band
MSG_PEEK Peek at data
MSG_DONTROUTE Bypass routing

- Send a UDP message on a socket to specified destination

```
#include <sys/socket.h>
```

```
int sendto (int sockfd, const void *buff ,int nbytes, int flags ,  
            const struct sockaddr *to , int addrlen ) ;
```

recvfrom System Calls

- receive a UDP message on a socket along with address of sending source

```
#include <sys/socket.h>
```

```
int recvfrom (int sockfd, const void *buff ,int nbytes, int flags ,  
              const struct sockaddr *from , int *addrlen ) ;
```

UDP Echo Sever

UDP Server
\$./udpserver Port

```
int main(int argc, char **argv)
{
    int      sockfd ,len;  char buff[ MAX];
    struct sockaddr_in  servaddr, cliaddr;

    sockfd = socket(AF_INET, SOCK_DGRAM, 0);

    bzero(&servaddr, sizeof(servaddr));

    servaddr.sin_family      = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port        = htons (atoi(argv[1]));

    bind(sockfd, (struct Sockaddr *) &servaddr, sizeof(servaddr));

    for (;;)
    {
        len = clilen;
        n=recvfrom(sockfd, buff, MAX, 0, cliaddr, &len);
        printf(" Received Msg.:%s\n",buff);

        sendto(sockfd, buff, n, 0, (struct sockaddr*)&cliaddr, len);
    }
}
```

UDP Echo Client

UDP Client

\$./udpserver IPADD Port

```
int main(int argc, char **argv)
{
    int      sockfd ,len; char buff[ MAX];
    struct sockaddr_in      servaddr;
    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);

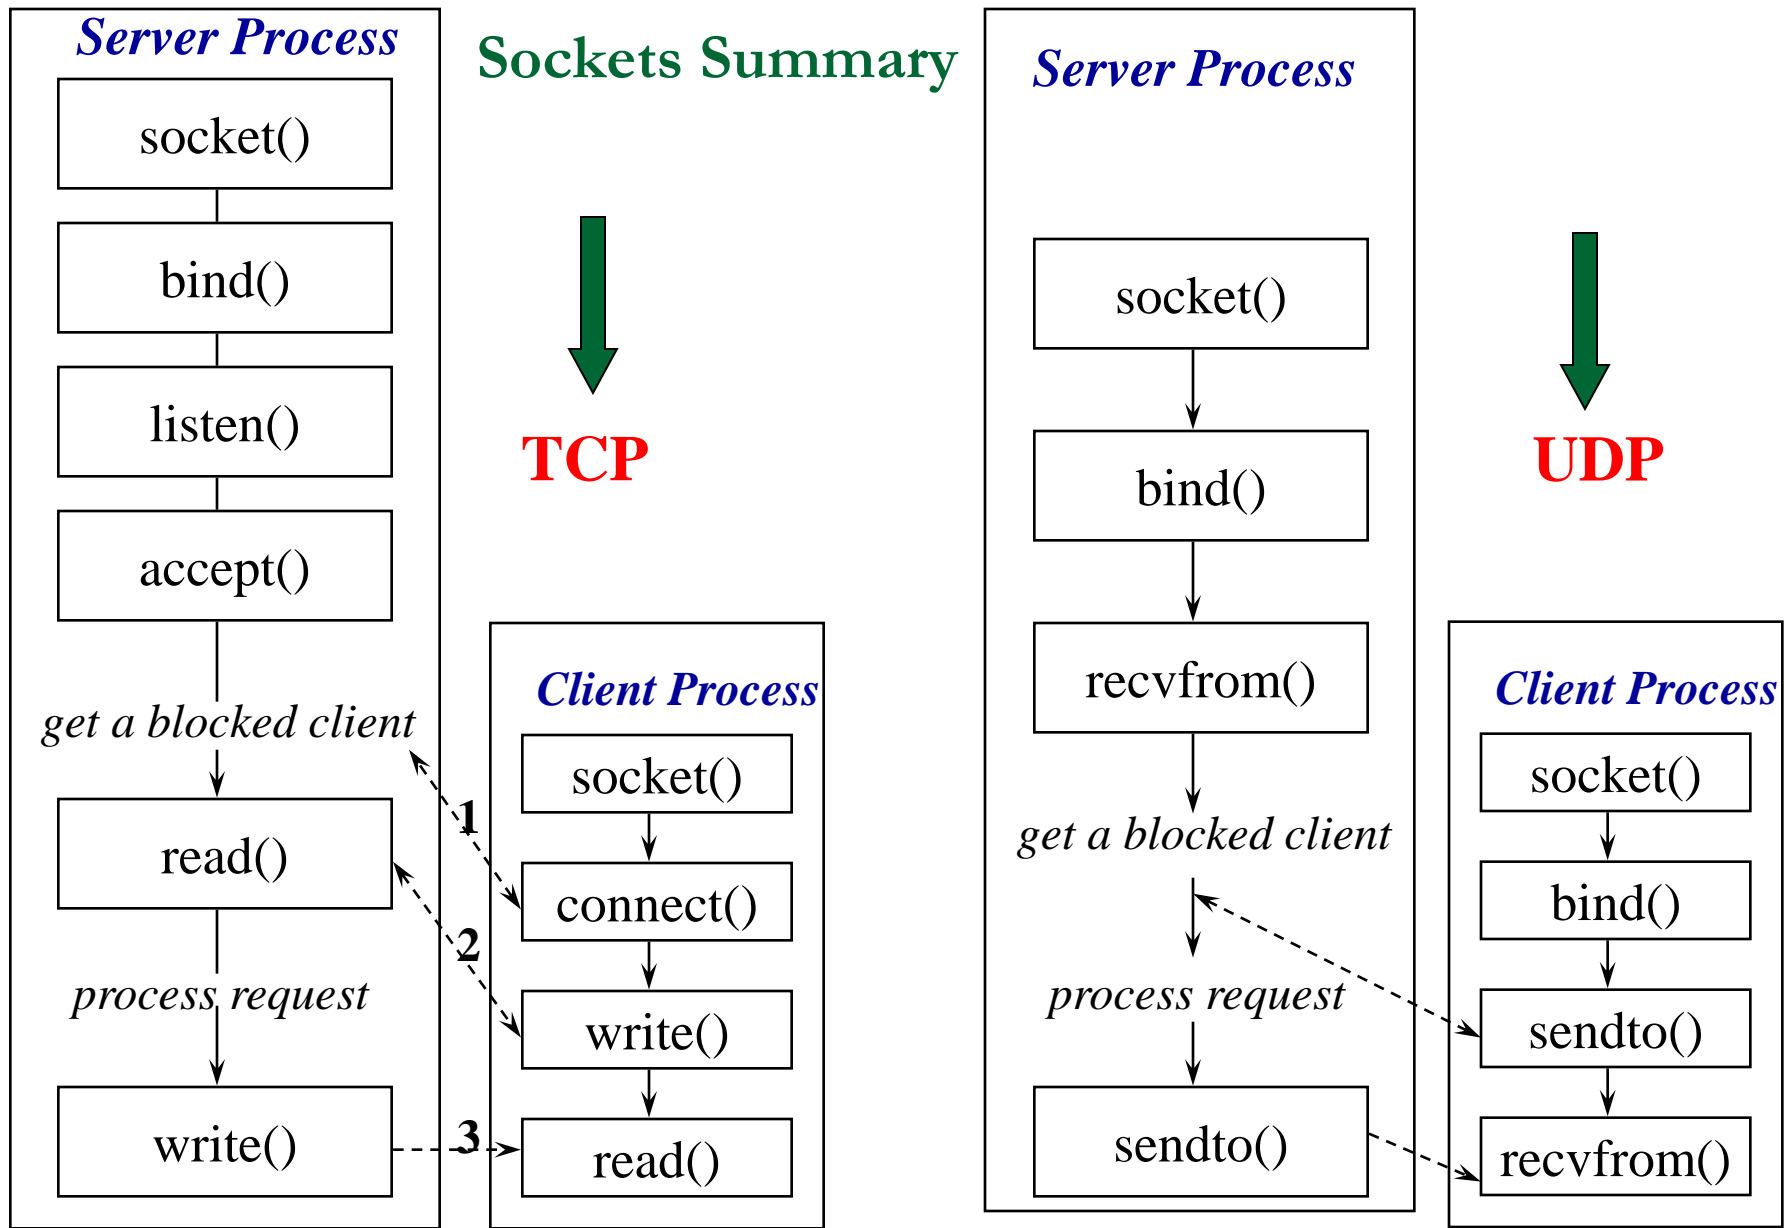
    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    inet_pton ( AF_INET, argv[1] , &servaddr.sin_addr );
    servaddr.sin_port = htons( atoi (argv[2]) );

    printf("Enter Text message\n");
    scanf( "%s",buff);
    sendto (sockfd, buff, MAX, 0,(struct sockaddr*)&servaddr, sizeof (servaddr) );

    recvfrom(sockfd, buff, MAX, 0, NULL, NULL);
    printf("Received Mesg.:%s\n",buff);

    exit(0);
}
```


Sockets Summary



```
#include <arpa/inet.h> #include <sys/socket.h> #include <netinet/in.h>
```

```
#define SERVER_PORT 898989
```

```
Main() {
```

```
struct sockaddr_in cliaddr,servaddr;
```

```
int sockfd,confd,clilen;
```

```
sockfd = socket(AF_INET , SOCK_STREAM , 0) ;
```

```
servaddr.sin_family = AF_INET;
```

```
servaddr.sin_port = htons(SERVER_PORT);
```

```
servaddr.sin_addr.s_addr =htonl( INADDR_ANY) ;
```

```
bind(sockfd, (struct sockaddr *)&servaddr,sizeof(servaddr));
```

```
listen(sockfd,10);
```

```
while(1) {
```

```
clilen=sizeof(cliaddr);
```

```
confd= accept(sockfd,(struct sockaddr *)&cliaddr,&clilen);
```

```
/* printf("Client IP: %s\n", inet_ntoa(cliaddr.sin_addr)); */
```

```
/* printf("Client Port: %hu\n", ntohs(cliaddr.sin_port)); */
```

```
read(confd, buf, BUFFER_SIZE); printf("Received:%s\n",buf);
```

```
write(confd, buf, BUFFER_SIZE);
```

```
close(confd);
```

```
}
```

```
return 0;
```

```
}
```

Iterative Server

\$/server

```
#include <arpa/inet.h> #include <sys/socket.h> #include <netinet/in.h>
```

```
#define SERVER_PORT 898989
```

```
main () {
```

```
    struct sockaddr_in cliaddr,servaddr;
```

```
    int sockfd,confd,clilen;
```

```
    sockfd = socket(AF_INET , SOCK_STREAM , 0) ;
```

```
    servaddr.sin_family = AF_INET;
```

```
    servaddr.sin_port = htons(SERVER_PORT);
```

```
    servaddr.sin_addr.s_addr =inet_addr(argv[1]) ; // or inet_pton()
```

```
    connect(sockfd, (struct sockaddr *) &servaddr, sizeof(servaddr)) ;
```

```
    clilen = sizeof(cliaddr);
```

```
    /* getsockname(sockfd, (struct sockaddr *) &cliaddr, &clilen); */
```

```
    /*    printf("Client socket has IP: %s\n", inet_ntoa(cliaddr.sin_addr)); */
```

```
    /*    printf("Client socket has Port: %hu\n", ntohs(cliaddr.sin_port)); */
```

```
    printf("Enter data:\n"); scanf("%s",buf) ;
```

```
    write(sockfd, buf, BUFFER_SIZE);
```

```
    read(sockfd, buf, BUFFER_SIZE);    printf("Received Data :%s\n",buf);
```

```
    close(sockfd);
```

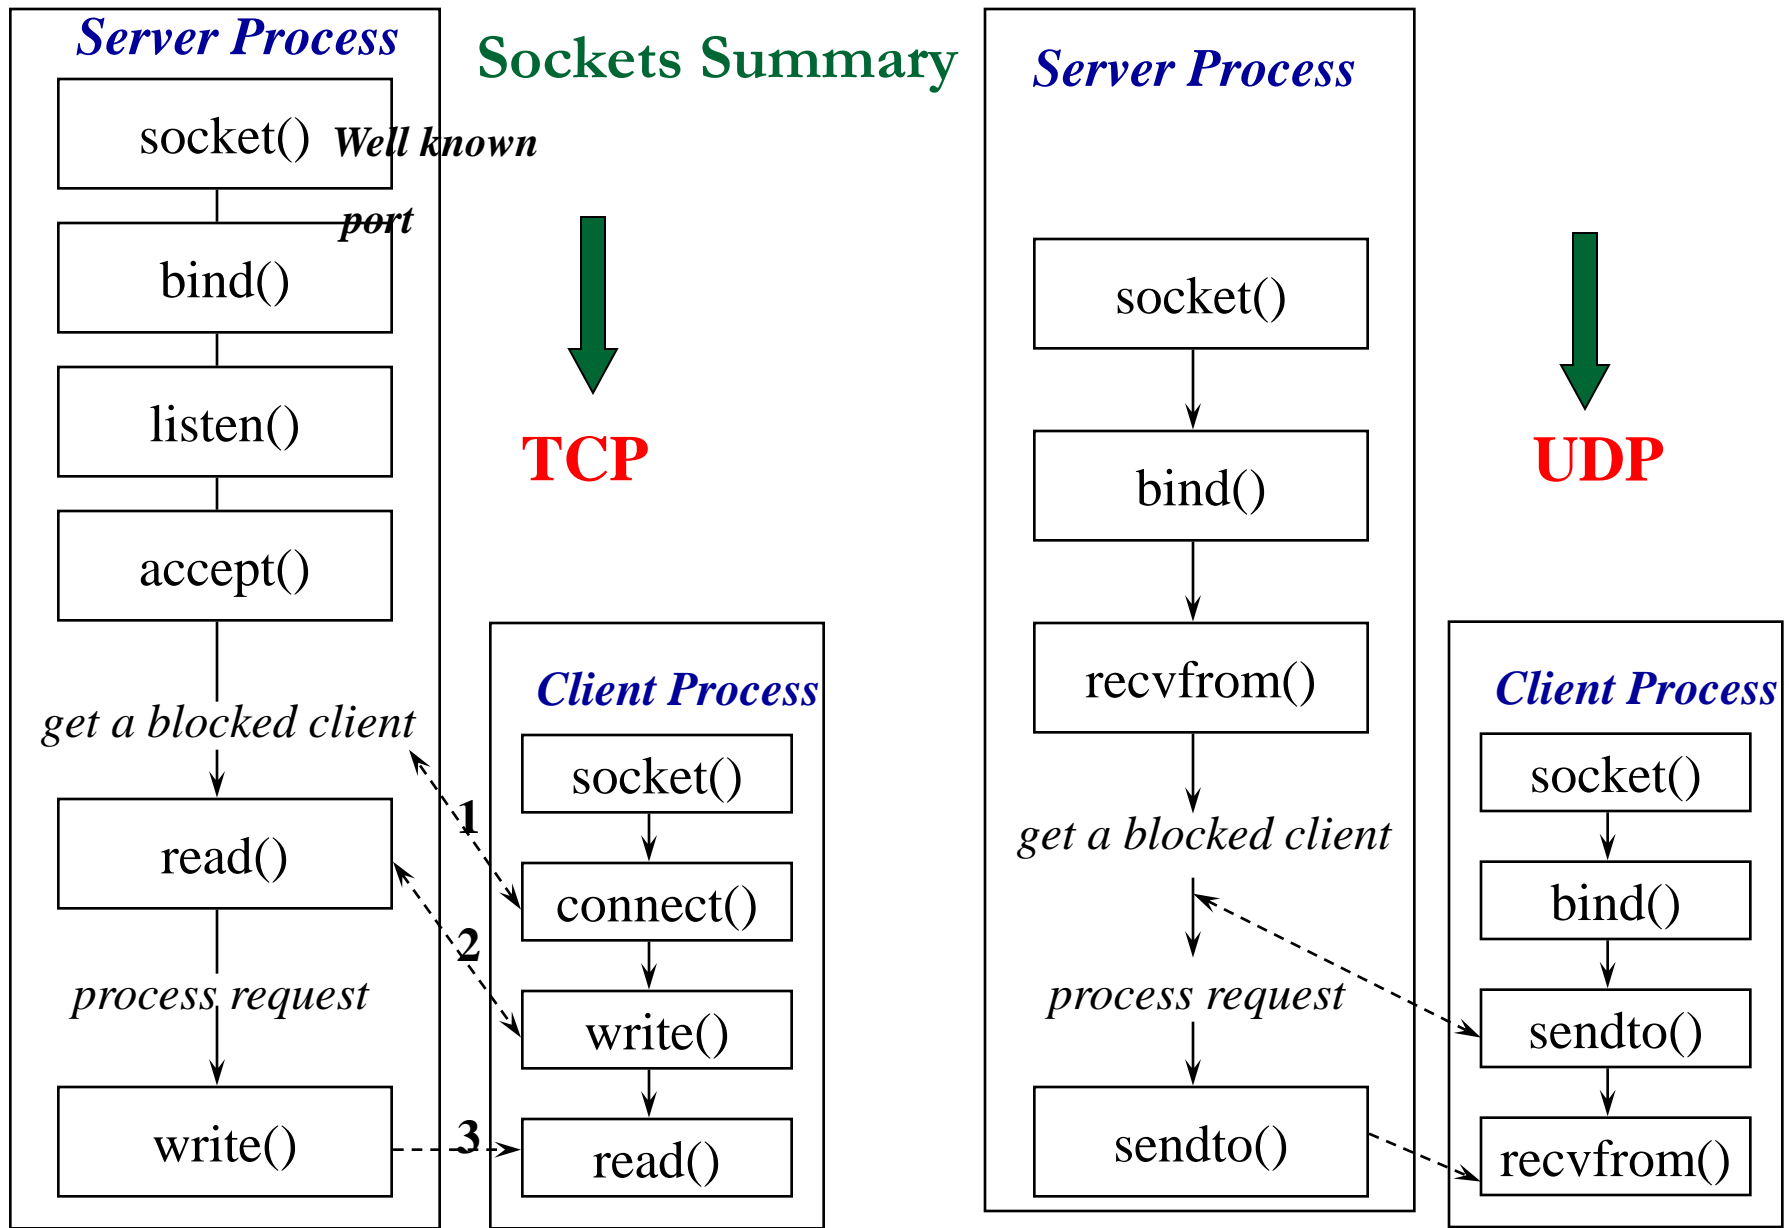
```
    exit(0);
```

```
}
```

Client Program

\$/client 172.24.2.4

Sockets Summary



UDP Echo Sever

UDP Server
\$./udpsvr Port

```
int main(int argc, char **argv)
{
    int      sockfd ,len;  char buff[ MAX];
    struct sockaddr_in  servaddr, cliaddr;

    sockfd = socket(AF_INET, SOCK_DGRAM, 0);

    bzero(&servaddr, sizeof(servaddr));

    servaddr.sin_family      = AF_INET;
    servaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    servaddr.sin_port        = htons (atoi(argv[1]));

    bind(sockfd, (struct Sockaddr *) &servaddr, sizeof(servaddr));

    for (;;)
    {
        len = clilen;
        n=recvfrom(sockfd, buff, MAX, 0, cliaddr, &len);
        printf(" Received Msg.:%s\n",buff);

        sendto(sockfd, buff, n, 0, (struct sockaddr*)&cliaddr, len);
    }
}
```

UDP Echo Client

UDP Client

\$./udpserver IPADD Port

```
int main(int argc, char **argv)
{
    int      sockfd ,len; char buff[ MAX];
    struct sockaddr_in      servaddr;
    sockfd = Socket(AF_INET, SOCK_DGRAM, 0);

    bzero(&servaddr, sizeof(servaddr));
    servaddr.sin_family = AF_INET;
    inet_pton ( AF_INET, argv[1] , &servaddr.sin_addr );
    servaddr.sin_port = htons( atoi (argv[2]) );

    printf("Enter Text message\n");
    scanf( "%s",buff);
    sendto (sockfd, buff, MAX, 0,(struct sockaddr*)&servaddr,  sizeof (servaddr) );
    recvfrom(sockfd, buff, MAX, 0, NULL, NULL);
    printf("Received Mesg.:%s\n",buff);
    exit(0);
}
```

Simple UDP Server

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#define SERVER_PORT    9988
int main()
{
    int    sockfd, clilen;
    char    buf[256];
    struct sockaddr_in servaddr, cliaddr;
sockfd = socket( AF_INET, SOCK_DGRAM, 0);
    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(SERVER_PORT);
    servaddr.sin_addr.s_addr =htonl(INADDR_ANY);
    if (bind(sockfd, (struct sockaddr*)&servaddr, sizeof(servaddr)) <0 )
        { printf("Server Bind Error"); exit(1); }
for( ; ; )
{ clilen= sizeof(cliaddr);
    recvfrom(sockfd,buf,256,0, (struct sockaddr*)&cliaddr,&clilen);

    printf("Server Received:%s\n",buf);

sendto(sockfd,"Server Got Message",18, 0, (struct
    sockaddr*)&cliaddr,clilen);
}
}
```

Simple UDP Client

```
#include <sys/types.h>
#include <sys/socket.h>
#include <netinet/in.h>
#define SERVER_PORT    9988
#define SERVER_IPADDR  "172.24.2.4"
int main()
{
    int    sockfd, len;
    char    buf[256];
    struct sockaddr_in , cliaddr, servaddr;

    servaddr.sin_family = AF_INET;
    servaddr.sin_port = htons(SERVER_PORT);
    servaddr.sin_addr.s_addr = inet_addr(SERVER_IPADDR);

    sockfd = socket( AF_INET, SOCK_DGRAM, 0);

    cliaddr.sin_family = AF_INET;
    cliaddr.sin_port = htons(0);
    cliaddr.sin_addr.s_addr = htonl(INADDR_ANY);
    bind(sockfd, (struct sockaddr*)&cliaddr, sizeof(cliaddr));

    printf("Enter Message\n");    fgets(buf, 255, stdin);
    len= sizeof(server);

    sendto(sockfd, buf, strlen(buf), 0, (struct sockaddr*)&servaddr, len);

    recvfrom(sockfd, buf, 256, 0, NULL, NULL);

    printf("Clinet Received: %s \n", buf);
    close(sockfd);
}
```

Not
mandatory

